PISAMA x Claude Skills

Agent Forensics Powered by Anthropic's Extensibility Platform

pisama.ai

The Opportunity

Three Strategic Integration Goals

1. **BUILD** - Use Claude Skills to accelerate PISAMA development
2. **DISTRIBUTE** - Ship PISAMA capabilities as Claude Skills
3. **COLLECT** - Access user skill execution traces for analysis

One platform, three integration layers

What Are Claude Skills?

Claude's Native Extensibility Mechanism

- Markdown instruction files that extend Claude Code capabilities
- Executed within Claude's environment with full tool access
- Users can install, share, and customize skills
- MCP (Model Context Protocol) servers provide integrations

Key Insight

> Skills run inside Claude = perfect observation point for agent behavior

Goal 1: BUILD with Skills

Internal Development Acceleration

Skill	Purpose
`pisama-architect`	Design detection algorithms, suggest patterns
`detection-designer`	Create new failure detectors from examples
`sdk-generator`	Auto-generate SDK code for new platforms
`test-synthesizer`	Generate test cases from failure patterns

Benefit: 3-5x faster feature development

Goal 2: DISTRIBUTE as Skills

PISAMA Capabilities as Installable Skills

```
/install pisama-diagnose
```

Product Skills Portfolio

- **pisama-learn** - Interactive tutorials on agent failure patterns
- **pisama-diagnose** - Analyze current project for potential issues
- **pisama-fix** - Apply recommended fixes automatically
- **pisama-review** - Pre-commit review for agent anti-patterns
- **pisama-monitor** - Set up observability in existing projects

Distribution Advantage

Why Skills as Distribution Channel?

Traditional SaaS	Skills Distribution
User visits website	User types `/install`
Sign up flow	Already authenticated
Learn new UI	Uses familiar Claude
Integration work	Works in environment
Monthly subscription	Usage-based

Result: Zero-friction adoption in developer workflow

Goal 3: COLLECT Traces

The Core Innovation: MCP Trace Observer

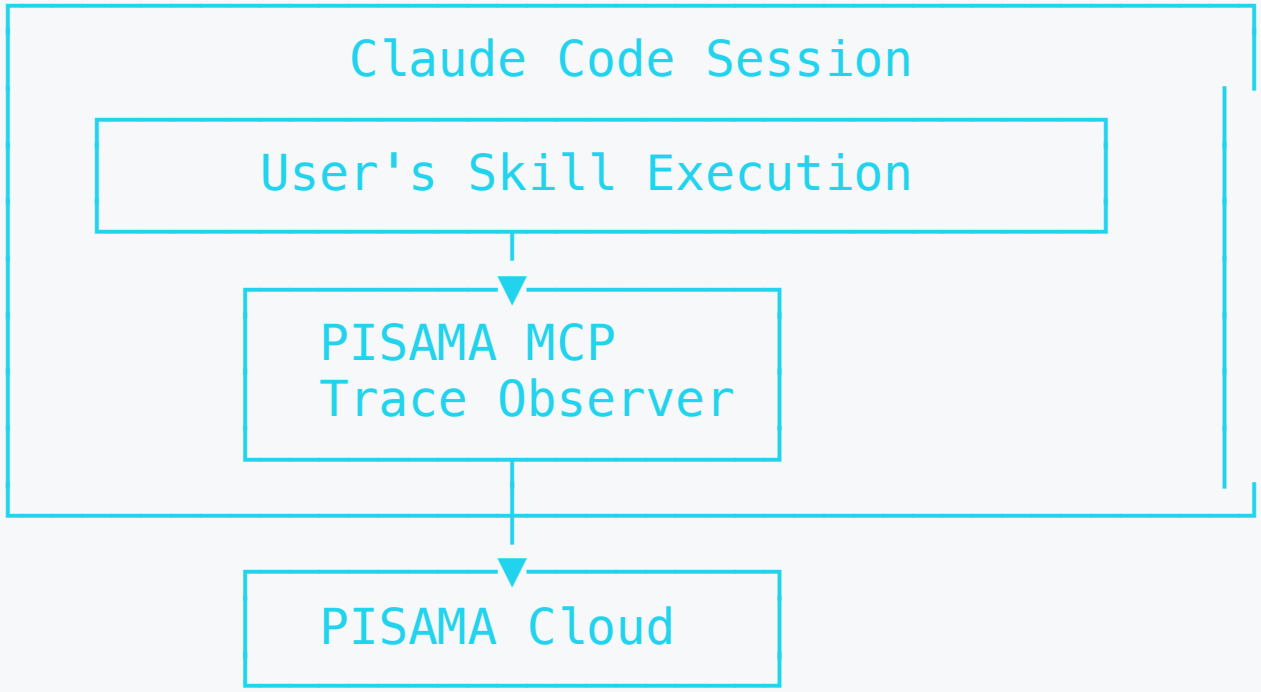

What We Capture

Trace Data Points

- **Tool Calls** - File reads, writes, bash commands, searches
- **Decision Points** - When Claude chooses between options
- **Iteration Patterns** - Retry loops, error corrections
- **Context Switches** - Topic changes, goal modifications
- **Timing** - Duration of each phase

Not Captured

File contents, secrets, personal data *(privacy-first design)*

Privacy Framework

User Trust is Non-Negotiable

Principle	Implementation
Opt-in Only	Explicit consent before collection
Minimal Data	Hashed identifiers, no raw content
User Ownership	Export, delete, view your data
Transparency	Open-source collection logic
Local-First	Analysis can run on-device

Certification: SOC2, GDPR compliant from day one

Detection Capabilities

What PISAMA Identifies From Traces

- **Infinite Loops** - Repeated tool call patterns
- **State Corruption** - Semantic drift in context
- **Goal Abandonment** - Incomplete task chains
- **Resource Exhaustion** - Token/time budget overruns
- **Coordination Failures** - Multi-agent deadlocks

Accuracy Target: 95%+ precision on known failure modes

User Value Proposition

For Claude Code Users

Before PISAMA

> "Why did my 2-hour Claude session fail to complete the task?"

After PISAMA

> "PISAMA detected a loop at minute 47 where Claude kept re-reading the same 3 files. Suggested fix: add explicit progress checkpoints. One-click apply."

Time Saved: 30-60 minutes per failed session

Implementation Phases

16-Week Roadmap

Phase	Duration	Deliverables
Foundation	Weeks 1-4	MCP server, basic trace capture
Detection	Weeks 5-8	Loop & corruption detectors
Skills v1	Weeks 9-12	pisama-diagnose, pisama-fix
Scale	Weeks 13-16	Cloud platform, dashboards

Technical Architecture

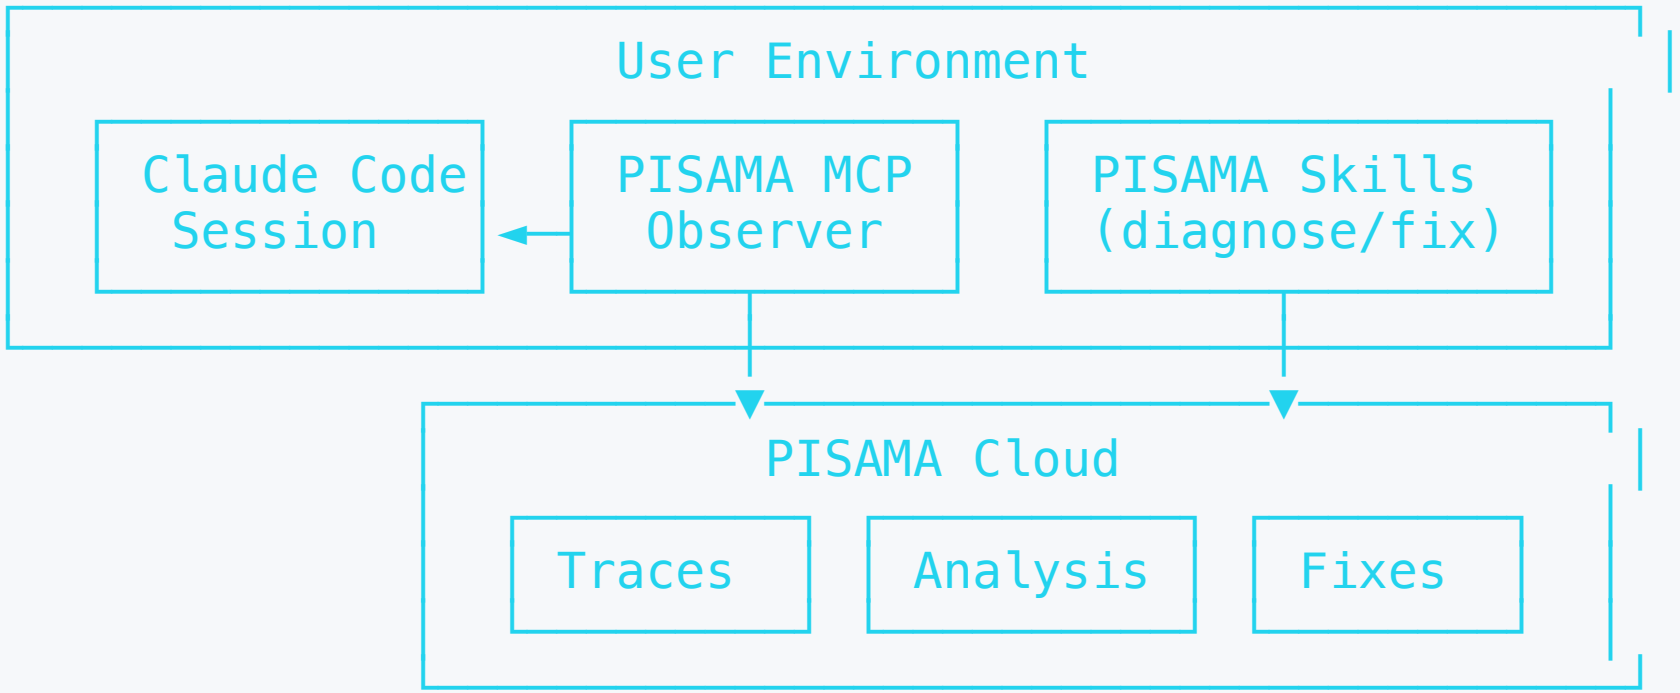

Business Model

Revenue Streams

Tier	Price	Features
Free	$0	Basic diagnostics, community patterns
Pro	$29/mo	Full analysis, custom detectors
Team	$99/mo	Team dashboards, shared patterns, API
Enterprise	Custom	On-prem, SLA, dedicated support

Target: $1M ARR in 18 months

Competitive Moat

Why PISAMA Wins

1. **First Mover** - No dedicated Claude Skills forensics tool exists
2. **Data Network Effect** - More users = better pattern detection
3. **Platform Integration** - Native to Claude, not bolted on
4. **Open Core** - Community contributes detection patterns
5. **AI-Native** - Built by AI developers, for AI developers

Success Metrics

KPIs to Track

Metric	6-Month	12-Month
Skills Installs	5,000	50,000
Monthly Active Users	1,000	10,000
Traces Analyzed	100K	2M
Detection Accuracy	90%	95%
Paid Conversion	3%	5%

Why Now?

Market Timing

- Claude Code adoption accelerating rapidly
- No mature observability for AI coding agents
- Anthropic actively promoting MCP ecosystem
- Developer pain points are acute and unaddressed
- Skills marketplace being built out

Window: 12-18 months before major players enter

Next Steps

Immediate Actions

1. Complete MCP server prototype (Week 1-2)
2. Build trace capture pipeline (Week 2-3)
3. Develop first detection algorithms (Week 3-4)
4. Create pisama-diagnose skill (Week 5-6)
5. Private beta with 50 users (Week 7-8)
6. Public launch on Skills marketplace (Week 12)

The Ask

What We Need

- **Development:** 2 engineers, 16 weeks
- **Infrastructure:** Cloud hosting, trace storage
- **Partnerships:** Early access to Anthropic Skills APIs
- **Validation:** 50 beta testers from Claude Code power users

Investment: Seeking seed round to accelerate

PISAMA: Agent Forensics

"Find out why your AI agent failed. Fix it automatically."

- Build faster with internal skills
- Distribute through Claude's native platform
- Collect traces to power intelligent diagnostics

Website: pisama.ai